Faster diameter computation in graphs of bounded Euler genus[*]

Kacper Kluk[†] Marcin Pilipczuk[‡] Michał Pilipczuk[§] Giannos Stamoulis[¶]

Abstract

We show that for any fixed integer $k \geqslant 0$, there exists an algorithm that computes the diameter and the eccentricies of all vertices of an input unweighted, undirected n-vertex graph of Euler genus at most k in time

$$\mathcal{O}_k(n^{2-\frac{1}{25}}).$$

Furthermore, for the more general class of graphs that can be constructed by clique-sums from graphs that are of Euler genus at most k after deletion of at most k vertices, we show an algorithm for the same task that achieves the running time bound

$$\mathcal{O}_k(n^{2-\frac{1}{356}} \log^{6k} n).$$

Up to today, the only known subquadratic algorithms for computing the diameter in those graph classes are that of [Ducoffe, Habib, Viennot; SICOMP 2022], [Le, Wulff-Nilsen; SODA 2024], and [Duraj, Konieczny, Potępa; ESA 2024]. These algorithms work in the more general setting of K_h-minor-free graphs, but the running time bound is $\mathcal{O}_h(n^{2-c_h})$ for some constant $c_h > 0$ depending on h. That is, our savings in the exponent, as compared to the naive quadratic algorithm, are independent of the parameter k.

The main technical ingredient of our work is an improved bound on the number of distance profiles, as defined in [Le, Wulff-Nilsen; SODA 2024], in graphs of bounded Euler genus.

1 Introduction

Computing the diameter of an input (undirected, unweighted) graph G is a classic computational problem that can be trivially solved in $\mathcal{O}(nm)$ time[1]. In 2013, Roditty and Vassilevska-Williams showed that this running time bound cannot be significantly improved in general: any algorithm distinguishing graphs of diameter 2 and 3 running in time $\mathcal{O}(m^{2-\varepsilon})$, for any fixed $\varepsilon > 0$, would break the Strong Exponential Time Hypothesis [14]. This motivates the search for restrictions on G that would make the problem of computing the diameter more tractable.

As shown by Cabello and Knauer [3], sophisticated orthogonal range query data structures allow near-linear diameter computation in graphs of constant treewidth. A breakthrough result by Cabello [2] showed that the diameter of an n-vertex planar graph can be computed in $\widetilde{\mathcal{O}}(n^{11/6})$ time; this complexity has been later improved by Gawrychowski, Kaplan, Mozes, Sharir, and Weimann to $\widetilde{\mathcal{O}}(n^{5/3})$ [8][2]. A subsequent line of research [5, 6, 11] generalized this result to K_h-minor-free graphs: for every integer h, there exists a constant $c_h > 0$ such that the diameter problem in n-vertex K_h-minor-free graphs can

[*]K.K. and Ma.P. are supported by Polish National Science Centre SONATA BIS-12 grant number 2022/46/E/ST6/00143. This work is a part of project BOBR (Mi.P., G.S.) that has received funding from the European Research Council (ERC) under the European Union's Horizon 2020 research and innovation programme (grant agreement No. 948057). In particular, a majority of work on this manuscript was done while G.S. was affiliated with University of Warsaw.

[†]Institute of Informatics, University of Warsaw, Poland. `k.kluk@uw.edu.pl`

[‡]Institute of Informatics, University of Warsaw, Poland. `m.pilipczuk@uw.edu.pl`

[§]Institute of Informatics, University of Warsaw, Poland. `michal.pilipczuk@mimuw.edu.pl`

[¶]IRIF, Université Paris Cité, CNRS, Paris, France. `giannos.stamoulis@irif.fr`

[1]We follow the convention that the vertex and the edge count of the input graph are denoted by n and m, respectively.

[2]The $\widetilde{\mathcal{O}}(\cdot)$ notation hides factors polylogarithmic in n, and the $\mathcal{O}_k(\cdot)$ notation hides factors depending on a parameter k.

be solved in time $\mathcal{O}_h(n^{2-c_h})$. In the works [6, 11], it holds that $c_h = \Omega\left(\frac{1}{h}\right)$; so the savings tend to zero as the size of the excluded clique minor increases.

However, known lower bounds, including the one of [14], does not exclude the possibility that c_h can be made a universal constant. That is, no known lower bound refutes the following conjecture:

Conjecture 1.1. *There exists a constant $c > 0$ such that, for every integer $h > 1$, the diameter problem in (unweighted, undirected) n-vertex K_h-minor-free graphs can be solved in time $\mathcal{O}_h(n^{2-c})$.*

Graphs of bounded Euler genus. Our main result is the verification of Conjecture 1.1 for graphs of bounded Euler genus. Furthermore, our algorithm computes also the eccentricities of all the vertices of the input graph G. Recall here that the eccentricity of a vertex $v \in V(G)$ is defined as $\mathrm{ecc}(v) \coloneqq \max_{u \in V(G)} \mathrm{dist}_G(u, v)$, where $\mathrm{dist}_G(\cdot,\cdot)$ is the distance metric in G.

Theorem 1.2. *For every integers $k \geqslant 1$, there exists an algorithm that, given an (unweighted, undirected) n-vertex graph G of Euler genus at most k, runs in time $\mathcal{O}_k(n^{2-\frac{1}{25}})$ and computes the diameter of G and the eccentricity of every vertex of G.*

We remark that in [2, Section 9], Cabello briefly speculated that his approach could be also generalized to graphs embeddable on surfaces of bounded genus. However, as noted in [2], this would require significant effort, as the technique works closely on the embedding and in surfaces of higher genus, additional topological hurdles arise. In contrast, in our proof of Theorem 1.2 the main ingredient is an improved combinatorial bound on the number of so-called *distance profiles* [11] in graphs of bounded Euler genus. This proof uses topology only very lightly, while the rest of the argument is rather standard and topology-free. All in all, we obtain a robust methodology of approaching the problem, which, as we will see, can be also used to attack Conjecture 1.1 to some extent.

To explain our bound on distance profiles, we need to recall several relevant definitions.

Let G be a graph, $R \subseteq V(G)$ be a subset of vertices, and $s_R \in R$ be a vertex in R. The *distance profile* of a vertex $u \in V(G)$ to R (relative to s_R) is the function $\mathrm{prof}_{R,s_R}[u] \colon R \to \mathbb{Z}$ defined as follows:

$$\mathrm{prof}_{R,s_R}[u](s) = \mathrm{dist}_G(u,s) - \mathrm{dist}_G(u,s_R) \qquad \text{for all } s \in R.$$

Note that provided R is connected[3], we have $\mathrm{prof}_{R,s_R}[u](s) \in \{-|R|, -|R|+1, \ldots, |R|-1, |R|\}$. In [11], Le and Wulff-Nilsen proved that if R is connected and G is K_h-minor-free, then the set system

$$\big\{\{(s,i) \in R \times \{-|R|,\ldots,|R|\} \mid i \leqslant \mathrm{prof}_{R,s_R}[u](s)\} \;:\; u \in V(G)\big\}$$

has VC dimension at most $h-1$. Hence, by applying the Sauer-Shelah Lemma we obtain that

Theorem 1.3 ([11])**.** *For every integer $h \geqslant 1$, K_h-minor-free graph G, connected set $R \subseteq V(G)$, and $s_R \in R$, there are at most $\mathcal{O}_h(|R|^{2h-2})$ different distance profiles to R relative to s_R.*

The VC dimension argument applied above inevitably leads to a bound with the exponent depending on h. We show that for graphs of bounded Euler genus, the bound of Theorem 1.3 can be improved to a polynomial of degree independent of the parameter.

Theorem 1.4. *For every integer $k \geqslant 1$, (unweighted, undirected) graph G of Euler genus at most k, connected set $R \subseteq V(G)$, and $s_R \in R$, the number of distance profiles to R relative to s_R is at most $\mathcal{O}_k(|R|^{12})$.*

The main idea behind the proof of Theorem 1.4 is the following simple observation: if P is a shortest path from some $u \in V(G)$ to s_R, then, as one walks along P from u to s_R, the distance profile of the current vertex to R can only (point-wise) increase. A slightly more technical modification of this argument works for shortest paths from $u \in V(G)$ to R. This allows us to reduce the case of bounded Euler genus graphs to the planar case by cutting along a constant number of shortest-to-R paths, and analysing how the distance profiles change during such a process.

[3]A subset of vertices R of a graph G is *connected* if the induced subgraph $G[R]$ is connected.

One could ask whether an improvement similar to that of Theorem 1.4 would be possible even in the generality of K_h-minor-free graphs. Unfortunately, it seems that Theorem 1.4 is the limit of such improvements. More precisely, the following simple example shows that the linear dependency on h in the exponent of the bound on the number of profiles is inevitable even in graphs of treewidth h (which are $K_{(h+1)^2}$-minor-free).

Let $0 < k \ll \ell$ be positive integers. Let R be a path of length ℓ and $v_1, \ldots, v_k$ be k equidistant points on R (i.e., the distance between v_i and v_{i+1} is at least $p := \lfloor \ell/(k-1) \rfloor$). For every vector $\mathbf{a} = (a_1, \ldots, a_k) \in \{\ell, \ldots, \ell + p\}^k$, construct a vertex $u(\mathbf{a})$ and, for every $i \in \{1, \ldots, k\}$, connect it with v_i using a path of length a_i. This finishes the construction of the graph G; see Figure 1 for an illustration. Note that G has treewidth at most $k+1$, because $G - \{v_1, \ldots, v_k\}$ is a forest. Furthermore, since the distance between consecutive vertices v_i is at least p, we have that $\mathrm{dist}_G(u(\mathbf{a}), v_i) = a_i$ for every vector $\mathbf{a}$ and $i \in \{1, \ldots, k\}$. Consequently, if we restrict to vectors $\mathbf{a}$ with $a_1 = \ell$, every vertex $u(\mathbf{a})$ has a different distance profile to R relative to v_1. Finally, note that there are $(p+1)^{k-1} \geqslant (\ell/(k-1))^{k-1} = \Omega_k(\ell^k)$ different vectors $\mathbf{a}$ with $a_1 = \ell$, giving that many different profiles.

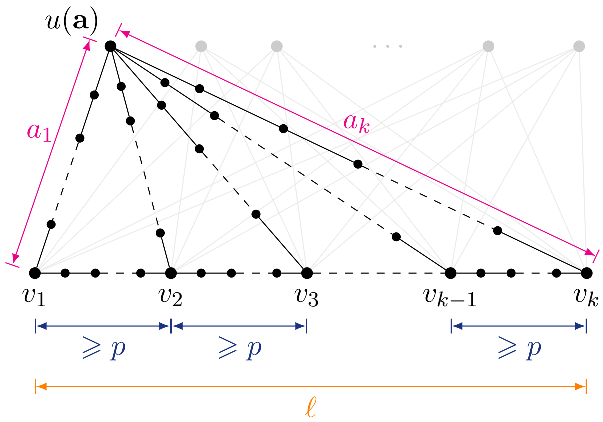

Figure 1: Illustration of a construction that shows that linear dependency on h in the exponent of the bound on the number of profiles is inevitable, even in graphs of treewidth h.

Our algorithm for Theorem 1.2 follows closely the approach of Le and Wulff-Nilsen [11] augmented by the bound provided by Theorem 1.4. Namely, we first compute an r-division of the input graph G into regions of size $r = n^\delta$, for some small $\delta > 0$. Then we use Theorem 1.4 for individual regions R to speed up the computation of distances between R and $V(G) - R$, by grouping vertices outside R according to their distance profiles to R. Each group is batch-processed in a single step.

Generalizations. Further, we show that our techniques combine well with the techniques for bounded treewidth graphs of Cabello and Knauer [3]. First, we show that Conjecture 1.1 holds for classes of graphs of bounded Euler genus with a constant number of *apices*, i.e., vertices that are arbitrarily connected to the rest of the graph.

Theorem 1.5. *For every integers $g, k \geqslant 1$, there exists an algorithm that, given an (unweighted, undirected) n-vertex graph G and a set $A \subseteq V(G)$ such that $|A| \leqslant k$ and $G - A$ is of Euler genus at most g, runs in time $\mathcal{O}_{g,k}(n^{2-\frac{1}{25}} \log^{k-1} n)$ and computes the diameter of G and the eccentricity of every vertex of G.*

Second, we show that Conjecture 1.1 holds for classes of graphs constructed by clique-sums of graphs as in Theorem 1.5. To state this result formally, we need some definitions. For a graph G, a *tree decompostion* of G is a pair (T, β) where T is a tree and β is a function that assigns to every $t \in V(T)$ a *bag* $\beta(t) \subseteq V(G)$ such that (1) for every $v \in V(G)$, the set $\{t \in V(T) \mid v \in \beta(t)\}$ is nonempty and connected in T, and (2) for every $uv \in E(G)$ there exists $t \in V(T)$ with $u, v \in \beta(t)$. The *torso* of the bag $\beta(t)$ is constructed from $G[\beta(t)]$ by adding, for every neighbor s of t in T, all edges between the vertices of $\beta(s) \cap \beta(t)$.

Theorem 1.6. *For every integer $k \geqslant 1$, there exists an algorithm with the following specification. The input consists of an (unweighted, undirected) n-vertex graph G together with a tree decomposition (T, β) of G and a set $A(t) \subseteq \beta(t)$ for every $t \in V(T)$ satisfying the following properties:*

- *For every node $t \in V(T)$, we have that $|A(t)| \leqslant k$ and the torso of $\beta(t)$ with the vertices of $A(t)$ deleted is a graph of Euler genus at most k.*
- *For every edge $st \in E(T)$, we have $|\beta(s) \cap \beta(t)| \leqslant k$.*

The algorithm runs in time $\mathcal{O}_k(n^{2-\frac{1}{356}} \log^{6k} n)$ and computes the diameter of G and the eccentricity of every vertex of G.

Note that the statements of Theorems 1.5 and 1.6 require the set A and the decomposition (T, β), respectively, to be provided explicitly on input; this should be compared with more general statements where the algorithm is given only G with a promise that such set A or decomposition (T, β) exist. At this point, we are not aware of any existing algorithm that would find in subquadratic time a set A as in Theorem 1.5, or the decomposition (T, β) with the sets A as in Theorem 1.6, even in the approximate sense. However, we were informed by Korhonen, Pilipczuk, Stamoulis, and Thilikos [10] that it seems likely that the techniques introduced in the recent almost linear-time algorithm for minor-testing [9] could be used to construct such an algorithm, with almost linear time complexity. With this result in place, the assumption about the decomposition and/or apex sets being provided on input could be lifted in Theorems 1.5 and 1.6; this is, however, left to future work.

Discussion. As one of the main outcomes of their theory of graph minors, Robertson and Seymour proved the following Structure Theorem [13]: every K_h-minor-free graph G admits a tree decomposition (T, β) such that

- for every pair s, t of adjacent nodes of T, the set $\beta(t) \cap \beta(s)$ has size $\mathcal{O}_h(1)$; and
- the torso of every bag $\beta(t)$ is "nearly embedable" into a surface of bounded (in terms of h) Euler genus.

The notion of being "nearly embeddable" encompasses adding a constant number of apices (which can be handled by Theorem 1.6) and a constant number of so-called vortices (which are not handled by Theorem 1.6). Thus, our methods fall short of verifying Conjecture 1.1 in full generality due to vortices.

We remark that recently, Thilikos and Wiederrecht [18] proved a variant of the Structure Theorem, where under the stronger assumption of excluding a minor of a *shallow vortex grid*, instead of a clique minor, they gave a decomposition as above, but with torsos devoid of vortices. Thus, the decomposition for shallow-vortex-grid-minor-free graphs provided by [18] can be directly plugged into Theorem 1.6, with the caveat that [18] does not provide a subquadratic algorithm to compute the decomposition.

Coming back to Conjecture 1.1, the simplest case that we are currently unable to solve is the setting when the input is a planar graph plus a single vortex. More formally, for a fixed integer k, let $\mathcal{G}_k$ be the class of graphs defined as follows. We have $G \in \mathcal{G}_k$ if there exist two subgraphs G_0, G_1 of G and a sequence of vertices $v_1, \ldots, v_b$ in $V(G_0) \cap V(G_1)$ such that:

- $V(G) = V(G_0) \cup V(G_1)$,
- $E(G) = E(G_0) \cup E(G_1)$,
- G_0 admits a planar embedding where the vertices $v_1, \ldots, v_b$ lie on one face in this order, and
- G_1 admits a tree decomposition (T_1, β_1), where T_1 is a path on nodes $t_1, \ldots, t_b$ and for every $i \in \{1, \ldots, b\}$, the bag $\beta_1(t_i)$ contains v_i and is of size at most k.

It is easy to see that graphs from $\mathcal{G}_k$ are $K_{k+\mathcal{O}(1)}$-minor-free. Do they satisfy Conjecture 1.1? That is, is there a constant $c > 0$ such that the diameter problem in $\mathcal{G}_k$ can be solved in time $\mathcal{O}_k(n^{2-c})$?

Organization. We prove Theorem 1.4 in Section 3. Theorem 1.5 is proven in Section 4; note that Theorem 1.2 follows from Theorem 1.5 for $k = 1$. Theorem 1.6 is proven in Section 5.

2 Preliminaries

Set systems and VC-dimension. A *set system* is a collection $\mathcal{F}$ of subsets of a given set A, which we call *ground set* of $\mathcal{F}$. We say that a subset $Y \subseteq A$ is *shattered by* $\mathcal{F}$ if $\{Y \cap S : S \in \mathcal{F}\} = 2^Y$, that is, the intersections of Y and the sets in $\mathcal{F}$ contain every subset of Y. The *VC-dimension* of a set system $\mathcal{F}$ with ground set A is the size of the largest subset $Y \subseteq A$ shattered by $\mathcal{F}$. The notion of VC-dimension was introduced by Vapnik and Chervonenkis [19].

We will use the following well-known Sauer-Shelah Lemma [15, 16], which gives a polynomial upper bound on the size of a set system of bounded VC-dimension.

Lemma 2.1 (Sauer-Shelah Lemma)**.** *Let $\mathcal{F}$ be a set system with ground set A. If the VC-dimension of $\mathcal{F}$ is at most k, then $|\mathcal{F}| = \mathcal{O}(|A|^k)$.*

Basic graph notation. All our graphs are undirected. For a graph G, the neighborhood of a vertex u is defined as $N_G(u) = \{v : uv \in E(G)\}$ and for $X \subseteq V(G)$ we have $N_G(X) = \bigcup_{u \in X} N_G(u) - X$.

The *length* of a path P, denoted $|P|$, is the number of edges of P. For two vertices u, v of a graph G, the *distance* between u and v, denoted $\mathrm{dist}_G(u, v)$, is defined as the minimum length of a path in G with endpoints u and v. For every $v \in V(G)$ and set $R \subseteq V(G)$, we set $\mathrm{dist}_G(v, R) \coloneqq \min\{\mathrm{dist}_G(v, y) \colon y \in R\}$. For vertices x, y appearing on a path P, by $P[x, y]$ we denote the subpath of P with endpoints x and y. The set of vertices traversed by a path P is denoted by $V(P)$. In all above notation, we sometimes drop the subscript if the graph is clear from the context.

For a nonnegative integer q, we use the shorthand $[q] \coloneqq \{1, \ldots, q\}$. For a vertex $v \in V(G)$ and a set $X \subseteq V(G)$, we define the X*-eccentricity* of v as $\mathrm{ecc}_X(v) \coloneqq \max_{x \in X} \mathrm{dist}(v, x)$. Thus, the eccentricity of v in G is the same as its $V(G)$-eccentricity.

The *Euler genus* of a graph G is the minimum Euler characteristic of a surface, where G is embeddable. We refer to the textbook of Mohar and Thomassen for more on surfaces and embedded graphs [12].

We will use the following result of Le and Wulff-Nilsen [11, Theorem 1.3] for planar graphs. Note that the set R is not necessarily connected.

Theorem 2.2. *Let $h \geqslant 1$ be an integer, G be a K_h-minor-free (unweighted, undirected) graph, R be a subset of $V(G)$, and $s_R \in R$. Then the set system*

$$\{\{(s, i) \in R \times \mathbb{Z} \mid i \leqslant \mathrm{dist}_G(u, s) - \mathrm{dist}_G(u, s_R)\} \ : \ u \in V(G)\}$$

has VC-dimension at most $h - 1$.

Algorithmic tools. All our algorithms assume the word RAM model.

To cope with apices, we will need the following classic data structure due to Willard [20].

Theorem 2.3 ([20])**.** *Let V be a set of n points in $\mathbb{R}^d$ and let $w \colon V \to \mathbb{R}$ be a weight function. By a* suffix range*, we mean any set of the form*

$$\mathsf{Range}(r_1, \ldots, r_d) \coloneqq \{(x_1, \ldots, x_d) \in \mathbb{R}^d \mid x_i \geqslant r_i \text{ for all } i \in [d]\}$$

for some range parameters $r_1, \ldots, r_d \in \mathbb{R}$.

There is a data structure that uses $\mathcal{O}\left(n \log^{d-1} n\right)$ preprocessing time, $\mathcal{O}\left(n \log^{d-1} n\right)$ memory and answers the following suffix range queries in time $\mathcal{O}\left(\log^{d-1} n\right)$: given a tuple $(r_i)_{i \in [d]}$, find the maximum value of $w(v)$ over all $v \in V \cap \mathsf{Range}(r_1, \ldots, r_d)$.

We will also need the following standard statement about r-divisions.

Theorem 2.4 ([21])**.** *Let G be a K_t-minor-free graph on n vertices. For any fixed constant $\varepsilon > 0$, and for any parameter r with $Ct^2 \log n \leqslant r \leqslant n$, where C is some absolute constant, we can construct in time $\mathcal{O}\left(n^{1+\varepsilon}\sqrt{r}\right)$ an r-division of G, that is, a collection $\mathcal{R}$ of connected subsets of vertices of G such that:*

- *$\bigcup \mathcal{R} = V(G)$,*
- *$|R| \leqslant r$ for every $R \in \mathcal{R}$, and*
- *$\sum_{R \in \mathcal{R}} |\partial R| \leqslant \mathcal{O}(nt/\sqrt{r})$, where $\partial R = R \cap N_G(V(G) - R)$.*

3 Distance profiles in graphs of bounded Euler genus

In this section we prove Theorem 1.4. Our argument consists of a reduction to the planar case, where we can use the constant bound on the VC-dimension of the set system given by the distance profiles due to Le and Wulff-Nilsen [11]. The main idea behind the reduction is to consider certain notions of "extended" profiles, where the extension is built along a collections of shortest paths. These shortest paths can be chosen in such a way that by cutting the graph along these paths we obtain a plane graph. Then a bound on the number of the extended profiles in the obtained plane graph translates to a bound on the number of (standard) distance profiles in the original graph.

Preliminary definitions and results needed for defining profiles with respect to shortest paths are given in Section 3.1. These extended profiles are then defined in Section 3.2. There, we also prove that a fundamental lemma that equality of extended profiles entails equality of (standard) distance profiles. The main reduction providing the proof of Theorem 1.4 is given at the end of this section.

3.1 Milestones

Let G be a graph, R be a subset of $V(G)$, v_0 be a vertex in $V(G)$, and P be a shortest path from v_0 to R. Let x be the unique vertex in $V(P) \cap R$. Further, let $\leqslant_P$ be the linear ordering of the vertices traversed by P: for two vertices $v, u \in V(P)$, we have $v \leqslant_P u$ if u belongs to $P[v, x]$. We say that a vertex $v \in V(P)$ is a *milestone of* P if either $v = x$ or we have $\mathrm{prof}_{R,x}[v] \neq \mathrm{prof}_{R,x}[u]$, where u is the successor of v in $\leqslant_P$. We denote by $M_R(P)$ the set of all milestones of P. Given a milestone $v \in M_R(P)$, the *neutral prefix of* v *in* P is defined as the vertex set of the maximal subpath Q of $P[v_0, v]$ satisfying the following: v is the only milestone of P that belongs to Q.

The next lemma shows that minimum-length paths towards R that contain a vertex in the neutral prefix of a milestone can be assumed to pass through that milestone vertex.

Lemma 3.1. *Let G be a graph, R be a subset of $V(G)$, v_0 be a vertex in $V(G)$ and P be a shortest path from v_0 to R. Then for every $v \in M_R(P)$, every u in the neutral prefix of v, and every $y \in R$, it holds that* $\mathrm{dist}(u, y) = |P[u, v]| + \mathrm{dist}(v, y)$.

Proof. TOPROVE 0 □

We also give an upper bound on the number of milestones.

Lemma 3.2. *Let G be a graph, R be a connected subset of $V(G)$, v_0 be a vertex of G, and P be a shortest path from v_0 to R. Then the number of milestones of P is at most $|R|^2 + 1$.*

Proof. TOPROVE 1 □

3.2 Anchor-distance profiles

Shortest path collections. Let G be a graph and R be a subset of vertices of G. We say that a collection $\mathcal{P}$ of paths in G is an *R-shortest path collection* if

- every $P \in \mathcal{P}$ is a shortest path from some $v^P \in V(G)$ to R, i.e., $|P| = \mathrm{dist}(v^P, R)$; and
- $R \subseteq \bigcup_{P \in \mathcal{P}} V(P)$.

For each $P \in \mathcal{P}$, we denote by x^P the endpoint of P in R. Note that the collection $\mathcal{P}$ obtained by taking, for every $y \in R$, the zero-length path from y to y, is an R-shortest path collection.

We say that an R-shortest path collection is *consistent* if, for every $P_1, P_2 \in \mathcal{P}$ and $v \in V(P_1) \cap V(P_2)$ the paths $P_1[v, x^{P_1}]$ and $P_2[v, x^{P_2}]$ are equal. That is, once two paths intersect, they continue together towards R.

The following statement is a direct consequence of the definition of an R-shortest path collection.

Observation 3.1. *Let G be a graph, R be a subset of vertices of G, and $\mathcal{P}$ be an R-shortest path collection. Then for every two paths $P_1, P_2 \in \mathcal{P}$ and every $v \in V(P_1) \cap V(P_2)$, we have $|P_1[v, x^{P_1}]| = |P_2[v, x^{P_2}]|$.*

Anchor vertices and their prefixes. Let G be a graph, R be a subset of $V(G)$, and $\mathcal{P}$ be an R-shortest path collection. We denote by $H_{\mathcal{P}}$ the union of the paths in $\mathcal{P}$, i.e., the graph $(\bigcup_{P\in\mathcal{P}} V(P), \bigcup_{P\in\mathcal{P}} E(P))$. We say that a vertex is an *anchor vertex* if either it has degree more than two in $H_{\mathcal{P}}$ or it is a milestone of a path $P \in \mathcal{P}$. We denote by $A_R(P)$ the set of all anchor vertices lying on a path $P \in \mathcal{P}$ and by $A_R(\mathcal{P})$ the set of all anchor vertices for $\mathcal{P}$. Given a path $P \in \mathcal{P}$ with endpoints v_0 and $y \in R$, and an anchor vertex $w \in A_R(P)$, the *prefix of* w *in* P is the vertex set of the maximal subpath Q of $P[v_0, v]$ satisfying the following: v is the only anchor vertex of P that belongs to Q. Note that for every anchor $w \in V(P)$ there is a milestone w' of P (possibly $w = w'$) such that the prefix of w in P is a subset of the neutral prefix of w' in P. Finally, for an anchor vertex w, the *tail* of w, denoted $\mathrm{tail}(w)$, is the subgraph of G consisting of the union of all prefixes of w in P over all paths $P \in \mathcal{P}$ that contain w.

Hat-distances. Let G be a graph, R be a subset of vertices of G, and $\mathcal{P}$ be an R-shortest path collection. We denote by

$$U_{\mathcal{P}} \coloneqq V(G) - \bigcup_{P\in\mathcal{P}} V(P).$$

For every $u \in U_{\mathcal{P}}$, and every anchor vertex $w \in A_R(\mathcal{P})$, we set

$$\widehat{\mathrm{dist}}(u, w) \coloneqq \min\{|Q_{u,z}| + |P[z, w]| \colon P \in \mathcal{P} \wedge w \in V(P) \wedge z \text{ is in the prefix of } w \text{ in } P\},$$

where $Q_{u,z}$ is a shortest path from u to z with all its internal vertices in $U_{\mathcal{P}}$. If such $Q_{u,z}$ does not exist for any $z \in V(\mathrm{tail}(w))$, we set $\widehat{\mathrm{dist}}(u, w) \coloneqq \infty$.

The following statement is a direct consequence of the definition of $\widehat{\mathrm{dist}}(\cdot, \cdot)$.

Observation 3.2. *Let G be a graph, R be a subset of vertices of G, and $\mathcal{P}$ be an R-shortest path collection. Then for every $u \in U_{\mathcal{P}}$, we have that*

$$\mathrm{dist}(u, R) = \min\left\{\widehat{\mathrm{dist}}(u, w) + \mathrm{dist}(w, R) \colon w \in A_R(\mathcal{P})\right\}.$$

Anchor-distance profiles. Let G be a graph, R be a subset of vertices of G, and $\mathcal{P}$ be an R-shortest path collection. The *anchor-distance profile* of a vertex $u \in U_{\mathcal{P}}$ to R with respect to $\mathcal{P}$ is a function $\mathrm{prof}^{\star}_{R,\mathcal{P}}[u]$ mapping each $w \in A_R(\mathcal{P})$ to

$$\mathrm{prof}^{\star}_{R,\mathcal{P}}[u](w) \coloneqq \widehat{\mathrm{dist}}(u, w) + \mathrm{dist}(w, R) - \mathrm{dist}(u, R).$$

Observation 3.2 implies that we always have $\mathrm{prof}^{\star}_{R,\mathcal{P}}[u](w) \geqslant 0$. We set

$$\widehat{\mathrm{prof}}_{R,\mathcal{P}}[u](w) \coloneqq \min\{\mathrm{prof}^{\star}_{R,\mathcal{P}}[u](w), |R| + 1\}.$$

Lemma 3.3. *Let G be a graph, let R be a connected subset of vertices of G, and $s_R \in R$. Also, let $\mathcal{P}$ be an R-shortest path collection. Then for all $u_1, u_2 \in U_{\mathcal{P}}$,*

$$\widehat{\mathrm{prof}}_{R,\mathcal{P}}[u_1] = \widehat{\mathrm{prof}}_{R,\mathcal{P}}[u_2] \qquad \textit{implies} \qquad \mathrm{prof}_{R,s_R}[u_1] = \mathrm{prof}_{R,s_R}[u_2].$$

Proof. TOPROVE 2 □

3.3 Reduction from bounded genus graphs to planar graphs

We next recall several definitions related to embeddings of graphs on surfaces. Our basic terminology follows [12]. We say that a graph H embedded in a surface Σ is a *simple cut-graph* of Σ if H has a single face that is also homeomorphic to an open disk; equivalently, H has a single facial walk. Given a surface Σ and a simple cut-graph H on Σ, we denote by $\Sigma \between H$ the surface obtained by cutting Σ along H. Note that, provided H is a simple cut-graph, $\Sigma \between H$ is always a disk.

Suppose now that a graph G embedded on Σ and H is a subgraph of G that is a simple cut-graph of H. We define $G \rtimes H$ to be the graph embedded on $\Sigma \rtimes H$ obtained from G as follows. First, let σ be the (unique) facial walk of H and note that each edge e of H is contained exactly twice in σ and each vertex v of H is contained in σ as many times as the degree of v in H. To obtain $G \rtimes H$, we replace H with a simple cycle C_σ whose vertex set is the set of copies of vertices of H and its edge set is the set of copies of edges of H in the obvious way. Notice that σ also prescribes for every edge uv of G between a vertex $u \in V(G) - V(H)$ and a vertex $v \in V(H)$, to which copy of v in $G \rtimes H$ the vertex u should be adjacent to (in $G \rtimes H$). See Figure 2 for an illustration.

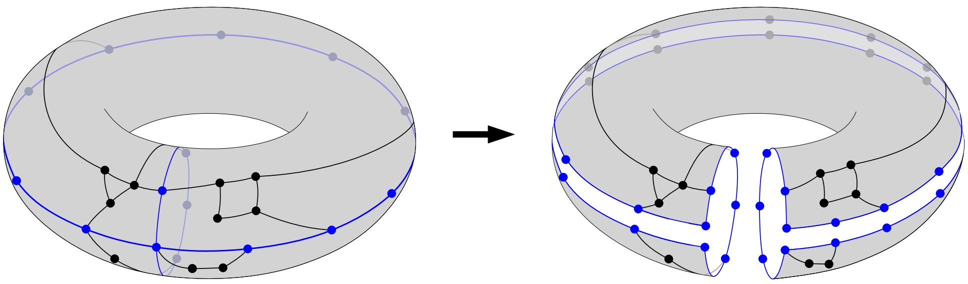

Figure 2: Left: A graph G embedded on a surface Σ and a subgraph H of G (in blue) that is a simple cut-graph of Σ. Right: The graph $G \rtimes H$ embedded on the surface $\Sigma \rtimes H$ (which is homeomorphic to a disk); the blue vertices/edges are copies of the vertices/edges of H.

We will use the following well-known result which appears in the literature under different formulations; see e.g. [1, 4, 7].

Lemma 3.4. *For every integer $k \geqslant 1$ and for every edge-weighted connected graph G embedded on a surface Σ of Euler genus at most k and every vertex $u \in V(G)$, there is a subgraph H of G with the following properties:*

- *H is a simple cut-graph of Σ, and*
- *$V(H)$ is the union of the vertex sets of $\mathcal{O}(k)$ shortest paths in G that have u as a common endpoint.*

We are now ready to proceed to the proof of Theorem 1.4. Employing Lemma 3.3, we aim at bouding the VC-dimension of the set system defined by the anchor-distance profiles. This is can be done by a suitable reduction to the planar setting using Lemma 3.4.

Proof. TOPROVE 3 $\square$

4 Bounded Euler genus graphs with apices: proof of Theorem 1.5

In this section we prove Theorem 1.5. (Note that Theorem 1.2 is a special case of Theorem 1.5 for $k = 1$.) We start by deriving the following corollary from Theorem 2.3.

Corollary 4.1. *Let V be a set of n points in $\mathbb{R}^d$. There is a data structure that uses $\mathcal{O}\left(dn \log^{d-2} n\right)$ preprocessing time, $\mathcal{O}\left(dn \log^{d-2} n\right)$ memory and answers the following queries in time $\mathcal{O}\left(d \log^{d-2} n\right)$: given $r_1, \ldots, r_d \in \mathbb{R}$, find $\max_{v \in V} \min_{i \in [d]} (v_i + r_i)$, where v_i denotes the ith coordinate of v.*

Proof. TOPROVE 4 $\square$

The main work in the proof of Theorem 1.5 will be done in the following lemma, which provides a fast computation of eccentricities once a suitable division is provided on input. We adopt the notation for divisions introduced in the statement of Theorem 2.4.

Lemma 4.2. *Fix constants $0 < \alpha, \gamma, \rho < 1$ and $k \in \mathbb{N}$. Assume we are given a connected graph G on n vertices with $O(n)$ edges with positive integer weights, a subset of vertices X, a subset of apices $A \subseteq V(G)$ of size at most k, and a family $\mathcal{R}$ with $V(G) - A = \bigcup \mathcal{R}$ such that the following conditions are satisfied:*

- $\sum_{R\in\mathcal{R}} |\partial R| \leqslant \mathcal{O}(n^\gamma)$;
- *for every* $R \in \mathcal{R}$, $|R| \leqslant \mathcal{O}(n^\rho)$ *and* $G[R]$ *is connected and contains* $O(|R|)$ *edges; and*
- *for every* $R \in \mathcal{R}$, *the number of distance profiles in* $G - A$ *on* ∂R *is of* $\mathcal{O}(n^\alpha)$.

Then, we can compute X*-eccentricity of every vertex of* G *in time* $\mathcal{O}(n^{\gamma+2\rho}\log n + (n^{1+\gamma}+n^{1+\alpha})\log^{k-1} n)$.

Proof. TOPROVE 5 □

The next statement is a reformulation of Theorem 1.5.

Theorem 4.3. *Fix constants* $k, g \in \mathbb{N}$. *Let* $\mathcal{C}$ *denote the class of all graphs that can be obtained by taking a graph* G *of Euler genus bounded by* g, *and adding* k *apices adjacent arbitrarily to the rest of* G *and to each other. Then there is an algorithm that given an unweighted graph* G *belonging to* $\mathcal{C}$, *together with its set of apices* A, *computes the eccentricity of every vertex in time* $\mathcal{O}_{k,g}\left(n^{1+\frac{24}{25}}\log^{k-1} n\right)$.

Proof. TOPROVE 6 □

5 The general case: Proof of Theorem 1.6

First, we show that data structure of Corollary 4.1 can be used to compute distances witnessed by shortest paths that pass through a constant-size separator.

Lemma 5.1. *Fix a constant* $k \in \mathbb{N}$. *There exists an algorithm which as the input receives an edge-weighted graph* G *on* n *vertices and* m *edges together with a partition of its vertices into three sets* A, B, C *such that* $|B| \leqslant k$ *and there are no edges between* A *and* C, *and as the output computes* $\max_{c\in C} \mathrm{dist}(a, c)$ *for every* $a \in A$. *The running time is* $\mathcal{O}(m \log n + n \log^{k-1} n)$.

Proof. TOPROVE 7 □

After computing the distances over a constant-size separator, we will use the following observation to simplify one of the sides of the separation.

Lemma 5.2. *Let* G *be a edge-weighted connected graph and let* A, B, C *be a partition of its vertices such that there are no edges between* A *and* C. *For every pair of vertices* $u, v \in B$, *let* $P_{u,v}$ *be any shortest path from* u *to* v *with all internal vertices in* C *(assuming such a path exists).*

Let G' *denote a graph obtained from* $G[A \cup B]$ *by adding an edge from* u *to* v *of weight equal to the length of* $P_{u,v}$, *for all* $u, v \in B$ *for which* $P_{u,v}$ *exists. Then,*

$$\mathrm{dist}_G(s, t) = \mathrm{dist}_{G'}(s, t) \qquad \textit{for all } s, t \in A \cup B.$$

Proof. TOPROVE 8 □

The next lemma encapsulates the main algorithmic content of the proof of Theorem 1.6. The algorithm will split the tree decomposition provided on input into smaller parts for which the eccentricities are easier to calculate. We use the following lemma to handle a single such part.

Lemma 5.3. *Fix constants* $k, g \in \mathbb{N}$, $0 < \delta < \frac{1}{54}$. *Assume we are given* $n \in \mathbb{N}$, *an edge-weighted graph* G *on at most* n *vertices with a weight function* $w\colon E(G) \to \mathbb{N}$, *a vertex subset* A *and a collection of non-empty vertex subsets* $V_0, V_1, \ldots, V_\ell$ *satisfying the following conditions:*

- *The sum of weights of all the edges in* G *is bounded by* $\mathcal{O}(n)$.
- $V(G) - A = V_0 \cup V_1 \cup \cdots \cup V_\ell$.
- $|A| \leqslant k$.
- *For every* $i \in [\ell]$, $G[V_i - V_0]$ *is connected,* $N_G(V_i - V_0) = V_i \cap V_0$, $|V_i| = \mathcal{O}(n^\delta)$, *and* $|V_0 \cap V_i| \leqslant 4$.
- *For all* $i, j \in [\ell], i \neq j$, $V_i - V_0$ *and* $V_j - V_0$ *are disjoint and non-adjacent in* G.
- *Every edge* $uv \in E(G)$ *with* $u, v \notin A$ *is contained in* $G[V_i]$ *for some* $i \in \{0, 1, \ldots, \ell\}$.

- *The graph obtained by taking $G[V_0]$ and adding a clique on $V_0 \cap V_i$ for every $i \in [\ell]$ has Euler genus bounded by g.*

Then, we can compute the eccentricity of every vertex of G in time $\mathcal{O}\left(n^{1+\frac{150+54\delta}{151}} \log^k n\right)$.

Proof. TOPROVE 9 □

Lemma 5.4. *Fix constants $k, g \in \mathbb{N}, 0 < \delta < \frac{1}{54}$. Assume we are given $n \in \mathbb{N}$, an edge-weighted graph G on at most n vertices with a weight function $w\colon E(G) \to \mathbb{N}$, a vertex subset A and a collection of non-empty vertex subsets $V_0, V_1, \ldots, V_\ell$ satisfying the same conditions as in Lemma 5.3 with the following differences:*

- *we don't require $G[V_i - V_0]$ to be connected and $V_i - V_0$ to be adjacent to whole $V_i \cap V_0$;*
- *instead of $|V_0 \cap V_i| \leqslant 4$, we require $|V_0 \cap V_i| \leqslant k$.*

Then, we can compute the eccentricity of every vertex of G in time $\mathcal{O}\left(n^{1+\frac{150+54\delta}{151}} \log^{k+5g} n\right)$.

Proof. TOPROVE 10 □

The next statement is a reformulation of Theorem 1.6.

Theorem 5.5. *Fix constants $k, g \in \mathbb{N}$. Assume we are given a graph G on n vertices together with its tree decomposition (T, β) and a set of private apices $A_t \subseteq \beta(t)$ for each node $t \in V(T)$ such that the following conditions hold:*

- *For every node $t \in V(T)$, we have $|A_t| \leqslant k$.*
- *For every edge $st \in E(T)$, we have $|\beta(v) \cap \beta(u)| \leqslant k$.*
- *For every node $t \in V(T)$, graph obtained by taking $G[\beta(t)] - A_t$ and turning $(\beta(t) \cap \beta(s)) - A_t$ into a clique for every edge $st \in E(T)$ has Euler genus bounded by g.*

Then, we can compute the eccentricity of every vertex of G in time $\mathcal{O}\left(n^{1+\frac{355}{356}} \log^{k+5g} n\right)$.

Proof. TOPROVE 11 □

Acknowledgements

Marcin thanks Jacob Holm, Eva Rotenberg, and Erik Jan van Leeuwen for many discussions on sub-quadratic algorithms for diameter while his stay on sabbatical in Copenhagen.

References

[1] Glencora Borradaile, Erik D. Demaine, and Siamak Tazari. Polynomial-time approximation schemes for subset-connectivity problems in bounded-genus graphs. *Algorithmica*, 68(2):287–311, 2014.

[2] Sergio Cabello. Subquadratic algorithms for the diameter and the sum of pairwise distances in planar graphs. *ACM Transactions on Algorithms*, 15(2):21:1–21:38, 2019.

[3] Sergio Cabello and Christian Knauer. Algorithms for graphs of bounded treewidth via orthogonal range searching. *Computational Geometry*, 42(9):815–824, 2009.

[4] Sergio Cabello, Éric Colin de Verdière, and Francis Lazarus. Algorithms for the edge-width of an embedded graph. *Computational Geometry*, 45(5):215–224, 2012. Special issue: 26th Annual Symposium on Computation Geometry at Snowbird, Utah, USA.

[5] Guillaume Ducoffe, Michel Habib, and Laurent Viennot. Diameter, eccentricities and distance oracle computations on H-minor free graphs and graphs of bounded (distance) Vapnik-Chervonenkis dimension. *SIAM Journal on Computing*, 51(5):1506–1534, 2022.

[6] Lech Duraj, Filip Konieczny, and Krzysztof Potępa. Better diameter algorithms for bounded VC-dimension graphs and geometric intersection graphs. In *32nd Annual European Symposium on Algorithms, ESA 2024*, volume 308 of *LIPIcs*, pages 51:1–51:18. Schloss Dagstuhl — Leibniz-Zentrum für Informatik, 2024.

[7] Jeff Erickson and Kim Whittlesey. Greedy optimal homotopy and homology generators. In *Proc. of the 16th Annual ACM-SIAM Symposium on Discrete Algorithms, SODA 2005*, pages 1038–1046. SIAM, 2005.

[8] Paweł Gawrychowski, Haim Kaplan, Shay Mozes, Micha Sharir, and Oren Weimann. Voronoi diagrams on planar graphs, and computing the diameter in deterministic $\widetilde{O}(n^{5/3})$ time. *SIAM Journal on Computing*, 50(2):509–554, 2021.

[9] Tuukka Korhonen, Michał Pilipczuk, and Giannos Stamoulis. Minor Containment and Disjoint Paths in almost-linear time. *CoRR*, abs/2404.03958, 2024.

[10] Tuukka Korhonen, Michał Pilipczuk, Giannos Stamoulis, and Dimitrios Thilikos, 2024. Private communication.

[11] Hung Le and Christian Wulff-Nilsen. VC set systems in minor-free (di)graphs and applications. In *2024 ACM-SIAM Symposium on Discrete Algorithms, SODA 2024*, pages 5332–5360. SIAM, 2024.

[12] Bojan Mohar and Carsten Thomassen. *Graphs on Surfaces*. Johns Hopkins series in the mathematical sciences. Johns Hopkins University Press, 2001.

[13] Neil Robertson and Paul D. Seymour. Graph Minors. XVI. Excluding a non-planar graph. *Journal of Combinatorial Theory, Series B*, 89(1):43–76, 2003.

[14] Liam Roditty and Virginia Vassilevska Williams. Fast approximation algorithms for the diameter and radius of sparse graphs. In *45th Symposium on Theory of Computing Conference, STOC 2013*, pages 515–524. ACM, 2013.

[15] Norbert Sauer. On the density of families of sets. *Journal of Combinatorial Theory, Series A*, 13(1):145–147, 1972.

[16] Saharon Shelah. A combinatorial problem; stability and order for models and theories in infinitary languages. *Pacific Journal of Mathematics*, 41(1):247 – 261, 1972.

[17] Saul Stahl and Lowell W. Beineke. Blocks and the nonorientable genus of graphs. *Journal of Graph Theory*, 1(1):75–78, 1977.

[18] Dimitrios M. Thilikos and Sebastian Wiederrecht. Killing a vortex. In *63rd IEEE Annual Symposium on Foundations of Computer Science, FOCS 2022*, pages 1069–1080. IEEE, 2022.

[19] V. N. Vapnik and A. Ya. Chervonenkis. On the uniform convergence of relative frequencies of events to their probabilities. *Theory of Probability & Its Applications*, 16(2):264–280, 1971.

[20] Dan E. Willard. New data structures for orthogonal range queries. *SIAM Journal on Computing*, 14(1):232–253, 1985.

[21] Christian Wulff-Nilsen. Separator theorems for minor-free and shallow minor-free graphs with applications. *CoRR*, abs/1107.1292, 2011.