

Light Spanners with Small Hop-Diameter

Sujoy Bhore* Lazar Milenković†

Abstract

Lightness, *sparsity*, and *hop-diameter* are the fundamental parameters of geometric spanners. Arya et al. [STOC'95] showed in their seminal work that there exists a construction of Euclidean $(1 + \varepsilon)$ -spanners with hop-diameter $O(\log n)$ and lightness $O(\log n)$. They also gave a general tradeoff of hop-diameter k and sparsity $O(\alpha_k(n))$, where α_k is a very slowly growing inverse of an Ackermann-style function. The former combination of logarithmic hop-diameter and lightness is optimal due to the lower bound by Dinitz et al. [FOCS'08]. Later, Elkin and Solomon [STOC'13] generalized the light spanner construction to doubling metrics and extended the tradeoff for more values of hop-diameter k . In a recent line of work [SoCG'22, SoCG'23], Le et al. proved that the aforementioned tradeoff between the hop-diameter and sparsity is tight for every choice of hop-diameter k . A fundamental question remains: *What is the optimal tradeoff between the hop-diameter and lightness for every value of k ?*

In this paper, we present a general framework for constructing light spanners with small hop-diameter. Our framework is based on *tree covers*. In particular, we show that if a metric admits a tree cover with γ trees, stretch t , and lightness L , then it also admits a t -spanner with hop-diameter k and lightness $O(kn^{2/k} \cdot \gamma L)$. Further, we note that the tradeoff for trees is tight due to a construction in uniform line metric, which is perhaps the simplest tree metric. As a direct consequence of this framework, we obtain a tight tradeoff between lightness and hop-diameter for doubling metrics in the entire regime of k .

1 Introduction

Let $M_X = (X, \delta_X)$ be a finite metric space, which can be viewed as a complete graph with vertex set X , where the weight of each edge $(u, v) \in \binom{X}{2}$ is equal to the metric distance between its endpoints, $\delta_X(u, v)$. Let $t \geq 1$ be a real parameter and let $H = (X, E)$ be a subgraph of M_X such that $E \subseteq \binom{X}{2}$. We say that H is a *t -spanner* for M_X , if for every two points u and v in X , it holds that $\delta_H(u, v) \leq t \cdot \delta_X(u, v)$, where $\delta_H(u, v)$ denotes the length of the shortest path between u and v in H . Such a path is called *t -spanner path* and parameter $t \geq 1$ is called the *stretch factor* (shortly, *stretch*) of H .

Spanners for Euclidean metric spaces (henceforth Euclidean spanners) are fundamental geometric structures with numerous applications, such as topology control in wireless networks [SVZ07], efficient regression in metric spaces [GKK17], approximate distance oracles [GLNS08], and more.

*Department of Computer Science & Engineering, Indian Institute of Technology Bombay, Mumbai, India.
Email: sujoy@cse.iitb.ac.in

†Tel-Aviv University. Email: milenkovic.lazar@gmail.com

Rao and Smith [RS98] showed the relevance of Euclidean spanners in the context of other geometric NP-hard problems, e.g., Euclidean traveling salesman problem and Euclidean minimum Steiner tree problem. Intensive ongoing research is dedicated to exploring diverse properties of Euclidean spanners; see [ADD⁺93, ADM⁺95, BKK⁺25, BT21a, BT21b, BT22, BCH⁺25, CG09, DNS95a, GLN02, KG92, ES15, RS98, LS19]. In fact, several distinct constructions have been developed for Euclidean spanners over the years, such as well-separated pair decomposition (WSPD) based spanners [Ca93, GLN02], skip-list spanners [AMS94], path-greedy and gap-greedy spanners [ADD⁺93, AS97], and many more; each such construction found further applications in various geometric optimization problems. For an excellent survey of results and techniques on Euclidean spanners, we refer to the book titled “Geometric Spanner Networks” by Narasimhan and Smid [NS07], and the references therein.

Besides having a small stretch, perhaps the most basic property of a spanner is its *sparsity*, defined as the number of edges in the spanner divided by $n - 1$, which is the size of an MST of the underlying n -point metric. Chew [Che86] was the first to show that there exists an Euclidean spanner with constant sparsity and stretch $\sqrt{10}$. Later, Keil and Gutwin [KG92] showed that the Delaunay triangulation is in fact a 2.42-spanner with constant sparsity. Clarkson [Cla87] designed the first Euclidean $(1 + \varepsilon)$ -spanner for \mathbb{R}^2 with sparsity $O(1/\varepsilon)$, for an arbitrary small $\varepsilon > 0$; an alternative algorithm was presented by Keil [Kei88]. These two papers [Cla87, Kei88] introduced the so-called Θ -graph as a new tool for designing $(1 + \varepsilon)$ -spanners with sparsity $O(1/\varepsilon)$ in \mathbb{R}^2 . Ruppert and Seidel [RS91] later generalized the Θ -graph to any constant dimension d , showing that one can construct a $(1 + \varepsilon)$ -spanner with sparsity $O(\varepsilon^{-d+1})$. Recently, Le and Solomon [LS19] showed that this bound is tight.

Another fundamental and extensively studied property of a spanner is its *lightness*, defined as the ratio of the sum of the edge weights of the spanner to the weight of the MST of the underlying metric. Das et al. [DHN93] showed that the “greedy spanner”, introduced by Althöfer et al. [ADD⁺93], has a constant lightness and stretch $1 + \varepsilon$ in \mathbb{R}^3 , for any constant $\varepsilon > 0$. This was generalized later by Das et al. [DNS95b] to \mathbb{R}^d , for all $d \in \mathbb{N}$. Rao and Smith [RS98] showed in their seminal work that the greedy spanner with stretch $1 + \varepsilon$ has lightness $\varepsilon^{-O(d)}$ in \mathbb{R}^d for every constant ε and d . After a long line of work, finally in 2019, Le and Solomon [LS19] improved the lightness bound of the greedy spanner to $O(\varepsilon^{-d} \log \varepsilon^{-1})$ in \mathbb{R}^d . For metrics with doubling dimension¹ d , Borradaile et al. [BLWN19] showed that the greedy spanner with stretch $1 + \varepsilon$ has lightness ε^{-2d} (see also [Got15] for an earlier work).

Besides having small *stretch* and *sparsity*, a spanner often possesses additional valuable properties of the underlying metric. One such critical property is the *hop-diameter*: a t -spanner for M_X has hop-diameter of k if, for any two points $u, v \in X$, there is a t -spanner path between u and v with at most k edges (or *hops*). Already in 1994, Arya et al. [AMS94] proposed a construction of Euclidean $(1 + \varepsilon)$ -spanners with logarithmic hop-diameter and a constant sparsity. In a subsequent work, Arya et al. [ADM⁺95] showed that there exists a construction of Euclidean $(1 + \varepsilon)$ -spanners with hop-diameter $O(\log n)$ and lightness $O(\log n)$. The same paper gives a general tradeoff of hop-diameter k and sparsity $O(\alpha_k(n))$ for $(1 + \varepsilon)$ -spanners, where α_k is an extremely slowly growing inverse of an Ackermann-style function (see also [BTS94, Sol13]). The former tradeoff of hop-diameter versus lightness is optimal due to the lower bound by Dinitz et al. [DES10]. Later, in

¹The doubling dimension of a metric space (X, δ_X) is the smallest value d such that every ball B in the metric space can be covered by at most 2^d balls of half the radius of B . A metric space is called doubling if its doubling dimension is constant.

2015, Elkin and Solomon [ES15] presented a light $(1+\varepsilon)$ spanner construction for doubling metrics and gave a tradeoff between the hop-diameter k and lightness for more values of k . Recent works by Le et al. [LMS22, LMS23] showed that the tradeoff of hop-diameter k and sparsity $O(\alpha_k(n))$ is asymptotically optimal, for every value of $k \geq 1$. However, despite a plethora of results on tradeoffs between *lightness* and *hop-diameter*, the following question remained open.

Question 1. *Given a set of points in \mathbb{R}^d , what is the optimal tradeoff between lightness and hop-diameter k , for every value of k ?*

A notion arguably stronger than that of a *spanner* is a *tree cover*. For a metric space $M_X = (X, \delta_X)$ let $T = (V_T, E_T)$ be a tree with $X \subseteq V_T$. We say that the tree T is *dominating*, if for every $u, v \in X$, it holds that $\delta_T(u, v) \geq \delta_X(u, v)$. A *tree cover* with *stretch* t is a collection of dominating trees such that for every pair of vertices $u, v \in X$, there exists a tree T in the collection with $\delta_T(u, v) \leq t \cdot \delta_X(u, v)$. The *size* of a tree cover is the number of trees in it. The lightness of a tree in the cover is the ratio of its weight to the weight of an MST of the underlying metric space. The *lightness* of a tree cover is the maximum lightness among the trees in the cover. We use L -light (γ, t) -tree cover to denote a tree cover with lightness L , γ trees, and stretch t . Clearly, the union of all the trees in a L -light (γ, t) -tree cover constitutes a t -spanner with sparsity bounded by $O(\gamma)$ and lightness bounded by $O(\gamma L)$.

The aforementioned tradeoff between hop-diameter and sparsity for Euclidean $(1+\varepsilon)$ -spanners by Arya et al. [ADM⁺95] was in fact achieved via tree covers. Their celebrated “*Dumbbell Theorem*” demonstrated that in \mathbb{R}^d , any set of points admits a tree cover of stretch $1+\varepsilon$ that uses only $O(\varepsilon^{-d} \cdot \log(1/\varepsilon))$ trees. Later, Bartal et al. [BFN22] generalized this theorem for doubling metrics. The bound on the tree cover size of [ADM⁺95] was recently improved by Chang et al. [CCL⁺24a] by a factor of $1/\varepsilon$. This is tight up to a logarithmic factor in $(1/\varepsilon)$ due to the lower bound on the sparse Euclidean spanners by Le and Solomon [LS19]. There are other tree cover constructions for doubling metrics [CGMZ16, BFN22], and for other metrics, such as planar and minor-free [BFN22, CCL⁺23, CCL⁺24b, GKR05, KLMN04] and general [BFN22, MN07, NT12].²

1.1 Our contributions

We present a general framework for constructing light spanners with small hop-diameter. Our starting point is the construction of a light 1-spanner with a bounded hop-diameter for tree metrics. The bounds are summarized in the following theorem, with the proof presented in Section 2.

Theorem 1.1. *For every $n \geq 1$, every $k \geq 1$, and every metric M_T induced by an n -vertex tree T , there is a 1-spanner for M_T with hop-diameter k and lightness $O(kn^{2/k})$.*

In order to go from tree metrics to arbitrary metrics, we rely on tree cover theorems. This reduction is summarized in the following corollary, with the proof provided in Section 4.

Corollary 1.2. *Let $n \geq 1$ be an arbitrary integer and let M_X be an arbitrary metric space with n points. If M_X admits an L -light (γ, t) -tree cover, then for any $k \geq 1$, the metric space M_X has a t -spanner with hop-diameter k and lightness $O(\gamma \cdot L \cdot k \cdot n^{2/k})$.*

²Metric induced by a graph $G = (V, E)$ is a metric space with point set V , where for every $u, v \in V$, their distance in the metric is equal to the shortest path distance between u and v in G , denoted by $\delta_G(u, v)$.

We note that the reduction from [Corollary 1.2](#) holds for any metric space which admits a light tree cover. To exemplify the reduction, we focus on doubling metrics and provide a general trade-off between hop-diameter and lightness for $(1 + \varepsilon)$ -spanners in the following theorem. (See [Section 4](#) for the proof.)

Theorem 1.3. *For every $k \geq 1$ and $\varepsilon \in (0, 1)$, every n -point metric space with doubling dimension d has a $(1 + \varepsilon)$ -spanner with hop-diameter k and lightness $O(\varepsilon^{-O(d)} \cdot kn^{2/k})$.*

Next, we compare our result with the aforementioned upper bound by Elkin and Solomon [[ES15](#)]. Since both constructions have a term $\varepsilon^{-O(d)}$ in lightness, we ignore those dependencies for clarity. The aforementioned construction of Elkin and Solomon [[ES15](#)] achieves hop-diameter $O(\log_\rho n + \alpha(\rho))$ and lightness $O(\rho \cdot \log_\rho n)$, for a parameter $\rho \geq 2$. In other words, the construction has hop-diameter $k' + O(\alpha(\rho))$ and lightness $k'n^{c/k'}$ for $k' \geq 1$, $\rho = n^{c/k'}$ and a constant c . Note that this tradeoff does not include values in the regime where hop-diameter is $o(\alpha(n))$. Namely, when $k' = O(\alpha(n))$, then $\rho = n^{\Omega(1/\alpha(n))}$ and the dominant term in hop-diameter is $\alpha(\rho) = \Theta(\alpha(n))$. In addition, the exponent $O(1/k)$ of n in the lightness is not asymptotically tight. Recall that the tradeoff we achieve is hop-diameter k versus lightness of $O(kn^{2/k})$, which holds for every value of $1 \leq k \leq n$. In other words, we give a fine-grained tradeoff for every value of k while nailing down the correct exponent of n , due to the lower bound we discuss next.

We complement our constructions with a lower bound for the n -point *uniform line metric*, which is a set of points on $[0, 1]$ with coordinates i/n for $0 \leq i \leq n - 1$. Refer to [Section 3](#) for the proof.

Theorem 1.4. *For every $n \geq 1$, every k such that $1 \leq k \leq n$, any spanner with hop-diameter k for the n -point uniform line metric must have lightness $\Omega(kn^{2/k})$.*

Note that, the lower bound holds for any value of stretch $t \geq 1$. Moreover, Dinitz et al. [[DES10](#)] previously obtained similar bounds, using a more intricate analysis based on linear programming method and connections to the so-called minimum linear arrangement problem. On the other hand, our proof is based on a rather simple combinatorial argument described in less than two pages. We believe such a simple combinatorial argument will have further implications in designing lower bounds for other related problems. The purpose of the lower bound is two-fold. First, since the uniform line metric is conceivably the simplest tree metric, we conclude that our upper bound for tree metrics ([Theorem 1.1](#)) is tight. In other words, our reduction ([Corollary 1.2](#)) is lossless. Second, the uniform line metric is also a doubling metric, which means that the tradeoff for the doubling metric ([Theorem 1.3](#)) of hop-diameter k and lightness $O(kn^{2/k})$ is tight.

2 Upper bound for tree metrics

This section is dedicated to proving [Theorem 1.1](#).

Theorem 1.1. *For every $n \geq 1$, every $k \geq 1$, and every metric M_T induced by an n -vertex tree T , there is a 1-spanner for M_T with hop-diameter k and lightness $O(kn^{2/k})$.*

We first describe the spanner construction in [Algorithm 1](#). The construction uses the following well-known result.

Lemma 2.1 ([[FL22](#)]). *Given any integer parameter $\ell > 0$ and an n -vertex tree T , there is a subset X of at most $\frac{2n}{\ell+1} - 1$ vertices such that every connected component of $T \setminus X$ has at most ℓ vertices and at most two outgoing edges towards X .*

Let $n \leftarrow |V(T)|$. If $n \leq k$, return the edge set of T . If $k = 1$, return the clique on $V(T)$. When $k = 2$, let c be the centroid of T , i.e., a vertex such that every component in $T \setminus \{c\}$ has size at most $n/2$. Let F be the returned set of edges, initialized to an empty set. For every vertex v in $V(T) \setminus \{c\}$, add to F the edge (c, v) of weight $\delta_T(c, v)$. Recurse with $k = 2$ on each of the subtrees of $T \setminus \{c\}$. When $k \geq 3$, let $\ell = k$ if $n \leq 2k^2$, and $\ell = \lfloor 2n^{2/k} \rfloor$ otherwise. Let X be the subset of V guaranteed by **Lemma 2.1** with parameter ℓ . Let T_1, \dots, T_g be the components of $T \setminus X$, each of which neighbors up to two vertices of X and each of which has size at most ℓ . For every connected component T_i of $T \setminus X$, let u_i and v_i be the vertices in X neighboring T_i . (It is possible that there is only one such vertex, but that case is handled analogously.) For each vertex w in T_i , add to F the edge (u_i, w) of weight $\delta_T(u_i, w)$ and (v_i, w) of weight $\delta_T(v_i, w)$. Recurse on T_i with parameter k . Consider a new tree T_X with a vertex set X and the edge set E_X consisting of: (1) all the edges in E with both endpoints in X and (line 25 in **Algorithm 1**); (2) edges (u_i, v_i) for every component T_i which has two neighbors in X (line 34 in **Algorithm 1**). Continue recursively for T_X and hop-diameter $k - 2$. This concludes the description of the algorithm.

Lemma 2.2 asserts that the constructed spanner has stretch 1 and hop-diameter k . **Lemma 2.3** asserts that the constructed spanner has lightness $O(kn^{2/k})$.

Lemma 2.2. *For every $k \geq 1, n \geq 1$ and every metric M_T induced by an n -vertex tree T , procedure $\text{SPANNER}(k, T)$ returns a 1-spanner of M_T with hop-diameter k .*

Proof. **TOPROVE 0** □

Lemma 2.3. *For every $k \geq 1, n \geq 1$ and every metric M_T induced by an n -vertex tree T of weight L , procedure $\text{SPANNER}(k, T)$ returns a spanner H_k of M_T with weight $W_k(n, L) = O(kn^{2/k}L)$.*

The lemma is implied by the following claims.

Claim 2.4. *For every $1 \leq n \leq k$, $W_k(n, L) \leq L$.*

Proof. **TOPROVE 1** □

Claim 2.5. *For every $n \geq 1$, $W_1(n, L) \leq \frac{n^2 L}{2}$.*

Proof. **TOPROVE 2** □

Claim 2.6. *For every $n \geq 1$, $W_2(n, L) \leq nL$.*

Proof. **TOPROVE 3** □

Claim 2.7. *Consider an invocation of $\text{SPANNER}(k, T)$ for $k \geq 3$ and let L_i be the weight of T_i plus the weight of the (at most two) edges connecting T_i to X . Then, $W_k(n, L) \leq 2\ell L + W_{k-2}(\frac{2n}{\ell+1}, L) + \sum_{i=1}^g W_k(\ell, L_i)$.*

Proof. **TOPROVE 4** □

Claim 2.8. *For every $n \geq 1$, $W_3(n, L) \leq 16n^{2/3}L$.*

Proof. **TOPROVE 5** □

Claim 2.9. *For every $1 \leq n \leq 16$, $W_k(n, L) \leq 9nL$.*

Proof. **TOPROVE 6** □

Claim 2.10. For every $1 \leq k < n \leq 8k$, $W_k(n, L) \leq 39kL$.

Proof. **TOPROVE 7** □

Claim 2.11. For every $k \geq 1$ and $1 \leq n \leq 2k^2$, $W_k(n, L) \leq 41kL$.

Proof. **TOPROVE 8** □

Claim 2.12. For all $k \geq 4$ and $n \geq 2k^2$, $W_k(n, L) \leq ckn^{2/k}$, for an absolute constant c .

Proof. **TOPROVE 9** □

3 Lower bound on the uniform line metric

This section is dedicated to proving **Theorem 1.4**, restated here for convenience.

Theorem 1.4. For every $n \geq 1$, every k such that $1 \leq k \leq n$, any spanner with hop-diameter k for the n -point uniform line metric must have lightness $\Omega(kn^{2/k})$.

Due to the inductive nature of the proofs, we consider a generalization of the uniform line metric.

Definition 3.1. Let $n \geq 1$ and $1 \leq p \leq n$ be arbitrary integers. A (p, n) line metric is a set of p points on $[0, 1]$ such that for every $0 \leq i \leq n-1$, the interval $[i/n, (i+1)/n)$ contains at most one point.

The uniform line metric with n points is an (n, n) line metric. The proof of **Theorem 1.4** follows from **Lemma 3.2** for $k = 1$, **Lemma 3.3** for $k = 2$, and **Lemma 3.4** for $k \geq 3$.

Lemma 3.2. For every n and p such that $1 \leq p \leq n$, let M be an arbitrary (p, n) line metric. Then, any spanner with hop-diameter 1 for M has lightness at least $W_1(p, n) \geq \frac{1}{64} \left(\frac{p^2}{n} \right)^2$.

Proof. **TOPROVE 10** □

Lemma 3.3. For every n and p such that $1 \leq p \leq n$, let M be an arbitrary (p, n) line metric. Then, any spanner with hop-diameter 2 for M has lightness at least $W_2(p, n) \geq \frac{1}{16} \cdot \frac{p^2}{n}$.

Proof. **TOPROVE 11** □

Lemma 3.4. For every n and p such that $1 \leq p \leq n$, let M be an arbitrary (p, n) line metric. Then, for any $k \geq 3$, any spanner with hop-diameter k for M has lightness at least $W_k(p, n) \geq ck \left(\frac{p^2}{n} \right)^{2/k}$, for $c = 1/73728$.

Proof. **TOPROVE 12** □




Figure 1: An illustration of the proof of the lower bound for $k = 2$ (Lemma 3.3). There is a vertex $p \in M_1$ (highlighted in red) which is not incident on any edge going to M_3 or M_4 . For every point $q \in M_4$, a 2-hop path from p to q induces a long edge, highlighted in red.

4 From light tree covers to light spanners

We start this section by proving Corollary 1.2, restated here for convenience.

Corollary 1.2. *Let $n \geq 1$ be an arbitrary integer and let M_X be an arbitrary metric space with n points. If M_X admits an L -light (γ, t) -tree cover, then for any $k \geq 1$, the metric space M_X has a t -spanner with hop-diameter k and lightness $O(\gamma \cdot L \cdot k \cdot n^{2/k})$.*

Proof. TOPROVE 13 □

Next, we use the following construction of light tree cover from [CCL⁺25].

Theorem 4.1 (Cf. Theorem 1.2 in [CCL⁺25]). *Given a point set P in a metric of constant doubling dimension d and any parameter $\varepsilon \in (0, 1)$, there exists an $\varepsilon^{-O(d)}$ -light $O(\varepsilon^{-O(d)}, 1 + \varepsilon)$ -tree cover for P .*

Corollary 1.2 and Theorem 4.1 immediately imply the following theorem.

Theorem 1.3. *For every $k \geq 1$ and $\varepsilon \in (0, 1)$, every n -point metric space with doubling dimension d has a $(1 + \varepsilon)$ -spanner with hop-diameter k and lightness $O(\varepsilon^{-O(d)} \cdot kn^{2/k})$.*

References

- [ADD⁺93] Ingo Althöfer, Gautam Das, David Dobkin, Deborah Joseph, and José Soares. On sparse spanners of weighted graphs. *Discrete & Computational Geometry*, 9(1):81–100, 1993. 2
- [ADM⁺95] Sunil Arya, Gautam Das, David M Mount, Jeffrey S Salowe, and Michiel Smid. Euclidean spanners: short, thin, and lanky. In *Proceedings of the twenty-seventh annual ACM symposium on Theory of computing*, pages 489–498, 1995. 2, 3
- [AMS94] Sunil Arya, David M Mount, and Michiel Smid. Randomized and deterministic algorithms for geometric spanners of small diameter. In *Proc. 35th IEEE Symposium on Foundations of Computer Science (FOCS)*, pages 703–712, 1994. 2
- [AS97] Sunil Arya and Michiel Smid. Efficient construction of a bounded-degree spanner with low weight. *Algorithmica*, 17(1):33–54, 1997. 2

- [BCH⁺25] Sujoy Bhore, Timothy M Chan, Zhengcheng Huang, Shakhar Smorodinsky, and Csaba D Toth. Sparse bounded-hop spanners for geometric intersection graphs. In *41st International Symposium on Computational Geometry, SoCG 2025, LIPIcs. Schloss Dagstuhl - Leibniz-Zentrum für Informatik*, 2025. 2
- [BFN22] Yair Bartal, Ora Nova Fandina, and Ofer Neiman. Covering metric spaces by few trees. *Journal of Computer and System Sciences*, 130:26–42, 2022. 3
- [BKK⁺25] Sujoy Bhore, Balázs Keszegh, Andrey Kupavskii, Hung Le, Alexandre Louvet, Dömötör Pálvölgyi, and Csaba D. Tóth. Spanners in planar domains via steiner spanners and non-steiner tree covers. In *Proceedings of the 2025 Annual ACM-SIAM Symposium on Discrete Algorithms, SODA 2025*, pages 4292–4326. SIAM, 2025. 2
- [BLWN19] Glencora Borradaile, Hung Le, and Christian Wulff-Nilsen. Greedy spanners are optimal in doubling metrics. In *Proc. 13th ACM-SIAM Symposium on Discrete Algorithms (SODA)*, pages 2371–2379, 2019. 2
- [BT21a] Sujoy Bhore and Csaba D. Tóth. Light euclidean steiner spanners in the plane. In *37th International Symposium on Computational Geometry, SoCG 2021, June 7-11, 2021, Buffalo, NY, USA (Virtual Conference)*, volume 189 of *LIPIcs*, pages 15:1–15:17, 2021. 2
- [BT21b] Sujoy Bhore and Csaba D. Tóth. On euclidean steiner $(1+\epsilon)$ -spanners. In *38th International Symposium on Theoretical Aspects of Computer Science, STACS 2021*, volume 187 of *LIPIcs*, pages 13:1–13:16, 2021. 2
- [BT22] Sujoy Bhore and Csaba D. Tóth. Euclidean steiner spanners: Light and sparse. *SIAM J. Discret. Math.*, 36(3):2411–2444, 2022. 2
- [BTS94] Hans L. Bodlaender, Gerard Tel, and Nicola Santoro. Trade-offs in non-reversing diameter. *Nord. J. Comput.*, 1(1):111–134, 1994. 2
- [Cal93] Paul B. Callahan. Optimal parallel all-nearest-neighbors using the well-separated pair decomposition. In *Proc. 34th IEEE Symposium on Foundations of Computer Science (FOCS)*, pages 332–340, 1993. 2
- [CCL⁺23] Hsien-Chih Chang, Jonathan Conroy, Hung Le, Lazar Milenkovic, Shay Solomon, and Cuong Than. Covering planar metrics (and beyond): $O(1)$ trees suffice. In *2023 IEEE 64th Annual Symposium on Foundations of Computer Science (FOCS)*, pages 2231–2261. IEEE, 2023. 3
- [CCL⁺24a] Hsien-Chih Chang, Jonathan Conroy, Hung Le, Lazar Milenkovic, Shay Solomon, and Cuong Than. Optimal euclidean tree covers. In *40th International Symposium on Computational Geometry, SoCG 2024, June 11-14, 2024, Athens, Greece*, volume 293 of *LIPIcs*, pages 37:1–37:15, 2024. 3
- [CCL⁺24b] Hsien-Chih Chang, Jonathan Conroy, Hung Le, Lazar Milenkovic, Shay Solomon, and Cuong Than. Shortcut partitions in minor-free graphs: Steiner point removal, distance oracles, tree covers, and more. In *SODA*, pages 5300–5331. SIAM, 2024. 3

- [CCL⁺25] Hsien-Chih Chang, Jonathan Conroy, Hung Le, Shay Solomon, and Cuong Than. Light tree covers, routing, and path-reporting oracles via spanning tree covers in doubling graphs. In *STOC*, 2025. 7
- [CG09] T-H Hubert Chan and Anupam Gupta. Small hop-diameter sparse spanners for doubling metrics. *Discrete & Computational Geometry*, 41:28–44, 2009. 2
- [CGMZ16] T-H Hubert Chan, Anupam Gupta, Bruce M Maggs, and Shuheng Zhou. On hierarchical routing in doubling metrics. *ACM Transactions on Algorithms (TALG)*, 12(4):1–22, 2016. 3
- [Che86] L. Paul Chew. There is a planar graph almost as good as the complete graph. In *Proc. 2nd Symposium on Computational Geometry*, pages 169–177. ACM Press, 1986. 2
- [Cla87] Kenneth L. Clarkson. Approximation algorithms for shortest path motion planning. In *Proc. 19th ACM Symposium on Theory of Computing (STOC)*, pages 56–65, 1987. 2
- [DES10] Yefim Dinitz, Michael Elkin, and Shay Solomon. Low-light trees, and tight lower bounds for euclidean spanners. *Discrete & Computational Geometry*, 43:736–783, 2010. 2, 4
- [DHN93] Gautam Das, Paul Heffernan, and Giri Narasimhan. Optimally sparse spanners in 3-dimensional Euclidean space. In *Proc. 9th Symposium on Computational Geometry (SoCG)*, pages 53–62. ACM Press, 1993. 2
- [DNS95a] Gautam Das, Giri Narasimhan, and Jeffrey Salowe. A new way to weigh malnourished euclidean graphs. In *proceedings of the sixth annual ACM-SIAM symposium on discrete algorithms*, pages 215–222, 1995. 2
- [DNS95b] Gautam Das, Giri Narasimhan, and Jeffrey S. Salowe. A new way to weigh malnourished Euclidean graphs. In *Proc. 6th ACM-SIAM Symposium on Discrete Algorithms (SODA)*, pages 215–222, 1995. 2
- [ES15] Michael Elkin and Shay Solomon. Optimal euclidean spanners: Really short, thin, and lanky. *J. ACM*, 62(5):35:1–35:45, 2015. 2, 3, 4
- [FL22] Arnold Filtser and Hung Le. Low treewidth embeddings of planar and minor-free metrics. In *FOCS*, pages 1081–1092. IEEE, 2022. 4
- [GKK17] Lee-Ad Gottlieb, Aryeh Kontorovich, and Robert Krauthgamer. Efficient regression in metric spaces via approximate Lipschitz extension. *IEEE Transactions on Information Theory*, 63(8):4838–4849, 2017. 1
- [GKR05] Anupam Gupta, Amit Kumar, and Rajeev Rastogi. Traveling with a pez dispenser (or, routing issues in mpls). *SIAM Journal on Computing*, 34(2):453–474, 2005. 3
- [GLN02] Joachim Gudmundsson, Christos Levkopoulos, and Giri Narasimhan. Fast greedy algorithms for constructing sparse geometric spanners. *SIAM Journal on Computing*, 31(5):1479–1500, 2002. 2

- [GLNS08] Joachim Gudmundsson, Christos Levcopoulos, Giri Narasimhan, and Michiel Smid. Approximate distance oracles for geometric spanners. *ACM Transactions on Algorithms (TALG)*, 4(1):1–34, 2008. 1
- [Got15] Lee-Ad Gottlieb. A light metric spanner. In *IEEE 56th Annual Symposium on Foundations of Computer Science, FOCS 2015, Berkeley, CA, USA, 17-20 October, 2015*, pages 759–772. IEEE Computer Society, 2015. 2
- [Kei88] J. Mark Keil. Approximating the complete Euclidean graph. In *Proc. 1st Scandinavian Workshop on Algorithm Theory (SWAT)*, volume 318 of *LNCS*, pages 208–213. Springer, 1988. 2
- [KG92] J. Mark Keil and Carl A. Gutwin. Classes of graphs which approximate the complete Euclidean graph. *Discrete & Computational Geometry*, 7:13–28, 1992. 2
- [KLMN04] Robert Krauthgamer, James R. Lee, Manor Mendel, and Assaf Naor. Measured descent: A new embedding method for finite metrics. In *FOCS*, pages 434–443. IEEE Computer Society, 2004. 3
- [LMS22] Hung Le, Lazar Milenkovic, and Shay Solomon. Sparse euclidean spanners with tiny diameter: A tight lower bound. In *38th International Symposium on Computational Geometry, SoCG 2022, June 7-10, 2022, Berlin, Germany*, volume 224 of *LIPIcs*, pages 54:1–54:15, 2022. 3
- [LMS23] Hung Le, Lazar Milenkovic, and Shay Solomon. Sparse euclidean spanners with optimal diameter: A general and robust lower bound via a concave inverse-ackermann function. In *39th International Symposium on Computational Geometry, SoCG 2023, June 12-15, 2023, Dallas, Texas, USA*, volume 258 of *LIPIcs*, pages 47:1–47:17, 2023. 3
- [LS19] Hung Le and Shay Solomon. Truly optimal Euclidean spanners. In *Proc. 60th IEEE Symposium on Foundations of Computer Science (FOCS)*, pages 1078–1100, 2019. 2, 3
- [MN07] Manor Mendel and Assaf Naor. Ramsey partitions and proximity data structures. *Journal of the European Mathematical Society*, 9(2):253–275, 2007. 3
- [NS07] Giri Narasimhan and Michiel Smid. *Geometric Spanner Networks*. Cambridge University Press, 2007. 2
- [NT12] Assaf Naor and Terence Tao. Scale-oblivious metric fragmentation and the nonlinear dvoretzky theorem. *Israel Journal of Mathematics*, 192:489–504, 2012. 3
- [RS91] Jim Ruppert and Raimund Seidel. Approximating the d -dimensional complete Euclidean graph. In *Proc. 3rd Canadian Conference on Computational Geometry (CCCG)*, pages 207–210, 1991. 2
- [RS98] Satish B. Rao and Warren D. Smith. Approximating geometrical graphs via “spanners” and “banyans”. In *Proc. 13th ACM Symposium on Theory of Computing (STOC)*, pages 540–550, 1998. 2
- [Sol13] Shay Solomon. Sparse euclidean spanners with tiny diameter. *ACM Trans. Algorithms*, 9(3):28:1–28:33, 2013. 2

- [SVZ07] Christian Schindelhauer, Klaus Volbert, and Martin Ziegler. Geometric spanners with applications in wireless networks. *Comput. Geom.*, 36(3):197–214, 2007. [1](#)

Algorithm 1: Procedure for constructing a spanner of a tree metric induced by a given tree $T = (V, E)$. Parameter $k \geq 1$ is the required hop-diameter. The procedure returns the edge set F of a spanner. The weight of every edge in F is assigned to be equal to the distance of its endpoints in T .

```

1: procedure SPANNER( $k, T = (V, E)$ )
2:    $n \leftarrow |V|$ 
3:   if  $n \leq k$  then
4:     return  $E$ 
5:   end if
6:   if  $k = 1$  then
7:     return  $\{(u, v) \mid u \in V, v \in V\}$  ▷ Clique on  $V$ 
8:   end if
9:   if  $k = 2$  then
10:    Let  $c$  be the centroid of  $T$ 
11:     $F \leftarrow \{(c, v) \mid v \in V \setminus \{c\}\}$  ▷ Connect  $c$  to every vertex in  $V \setminus \{c\}$ 
12:    Let  $T_1, \dots, T_g$  be the components of  $T \setminus \{c\}$ 
13:    for  $1 \leq i \leq g$  do
14:       $n_i \leftarrow |V(T_i)|$ 
15:       $F \leftarrow F \cup \text{SPANNER}(2, T_i)$ 
16:    end for
17:    return  $F$ 
18:   end if
19:    $\ell \leftarrow \lfloor 2n^{2/k} \rfloor$ 
20:   if  $k = 3$  then  $\ell \leftarrow \lfloor n^{2/3} \rfloor$ 
21:   end if
22:   if  $k \geq 4$  and  $n \leq 2k^2$  then  $\ell \leftarrow k$ 
23:   end if
24:   Let  $X$  be the set as in Lemma 2.1 with parameter  $\ell$  and let  $T_1, \dots, T_g$  be the components of  $T \setminus X$ 
25:    $E_X \leftarrow \{(u, v) \mid (u, v) \in E, u \in X, v \in X\}$ 
26:   for  $1 \leq i \leq g$  do
27:      $n_i \leftarrow |V(T_i)|$ 
28:      $F \leftarrow F \cup \text{SPANNER}(k, T_i)$ 
29:     if  $T_i$  has one neighbor  $u_i$  in  $X$  then  $F \leftarrow F \cup \{(u_i, w) \mid w \in V(T_i)\}$ 
30:     else
31:       Let  $u_i$  and  $v_i$  be the two neighbors of  $T_i$  in  $X$ 
32:        $F \leftarrow F \cup \{(u_i, w) \mid w \in V(T_i)\}$ 
33:        $F \leftarrow F \cup \{(v_i, w) \mid w \in V(T_i)\}$ 
34:        $E_X \leftarrow (u_i, v_i)$ 
35:     end if
36:   end for
37:    $F \leftarrow F \cup \text{SPANNER}(k-2, T_X = (X, E_X))$ 
38:   return  $F$ 
39: end procedure

```
